

Shaun McThomas
Thomas Hart
Project 3
Performance Tuning

For Parsing and entering batch data we performed the following optimizations:

Batch inserts: Rather than submitting and committing one record at a time, we perform batch inserts. This saves us preprocessing time on the database. This optimization was implemented from the start; therefore, it was not possible to get the actual time improvements over not using batch inserts.

Auto commit off: With auto commit turned off, we were able to reduce the overhead that the database introduces when writing to the disk with each insert.

Caching: The new values inserted the database - such as movies, stars, and genres - all contained values, such as the generated unique ID's, that were needed in further execution of the code. Rather than querying the database for these values, we returned them upon insertion and were able to keep them local for later reference.